1 A tiny word counter

We will build a minimal word counter and explain its parts. The program runs in $O(n)$ over the bytes of the input. We assemble the final script in the chunk `<<wc.py>>=` by referring to named subchunks that we define later.

⟨wc.py⟩≡

1.1 Imports

⟨imports⟩≡

```python
import sys
from collections import Counter
from pathlib import Path
import argparse
```

1.2 Argument parsing

⟨parse-args⟩≡

```python
def parse_args(argv=None):
    p = argparse.ArgumentParser(
        description="Count words in a file or stdin."
    )
    p.add_argument("file", nargs="?", help="Path to text file; default: stdin")
    p.add_argument("-n", "--top", type=int, default=10,
                   help="Show top-N words (default: 10)")
    return p.parse_args(argv)
```

1.3 Counting logic

⟨count-words⟩≡

```python
def words_from_text(text: str):
    # Very naive tokenization: split on whitespace and punctuation
    # Lowercasing makes counts case-insensitive.
    import re
    for w in re.findall(r"[A-Za-z0-9']+", text.lower()):
        yield w

def count_words(stream) -> Counter:
    c = Counter()
    for line in stream:
        c.update(words_from_text(line))
    return c
```

1.4 Program entry point

⟨main-guard⟩≡

```python
def main(argv=None):
    args = parse_args(argv)
```

```python
    if args.file:
        data = Path(args.file).read_text(encoding="utf-8", errors="ignore")
        stream = data.splitlines(keepends=True)
    else:
        stream = sys.stdin

    counts = count_words(stream)
    for word, n in counts.most_common(args.top):
        print(f"{n:7d}  {word}")

if __name__ == "__main__":
    main()
```