

A tiny word counter

We will build a minimal word counter and explain its parts.

The program runs in $O(n)$ over the bytes of the input.

We assemble the final script in the chunk <<wc.py>>= by referring to named subchunks that we define later.

wc.py

Imports

imports

```
import sys

from collections import Counter

from pathlib import Path

import argparse
```

Argument parsing

parse-args

```
def parse_args(argv=None):

    p = argparse.ArgumentParser(
        description="Count words in a file or stdin."
    )

    p.add_argument("file", nargs=?, help="Path to text file; default: stdin")

    p.add_argument("-n", "--top", type=int, default=10,
        help="Show top-N words (default: 10)")

    return p.parse_args(argv)
```

Counting logic

count-words

```
def words_from_text(text: str):
```

```

# Very naive tokenization: split on whitespace and punctuation

# Lowercasing makes counts case-insensitive.

import re

for w in re.findall(r"[A-Za-z0-9']+", text.lower()):
    yield w


def count_words(stream) -> Counter:
    c = Counter()

    for line in stream:
        c.update(words_from_text(line))

    return c

```

Program entry point

main-guard

```

def main(argv=None):

    args = parse_args(argv)

    if args.file:

        data = Path(args.file).read_text(encoding="utf-8", errors="ignore")

        stream = data.splitlines(keepends=True)

    else:

        stream = sys.stdin

    counts = count_words(stream)

    for word, n in counts.most_common(args.top):

        print(f"{n:7d} {word}")

if __name__ == "__main__":
    main()

```